

(IM) PRACTICAL FUNCTIONAL PROGRAMMING

ADOPTING FP IN INDUSTRY

WHY THIS TALK?

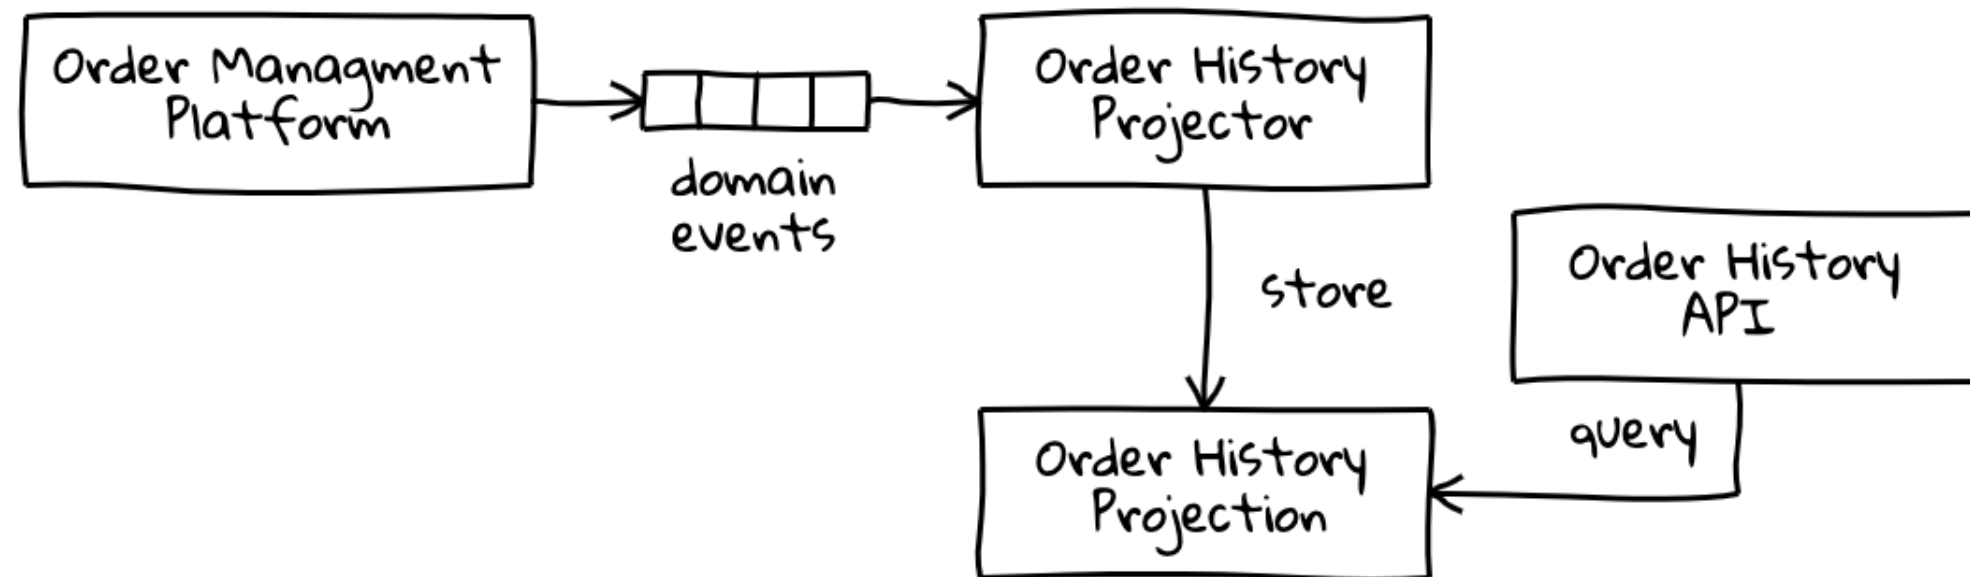
HOW MANY TIMES HAVE YOU HEARD:

- ▶ **FP is** too hard
- ▶ **FP is** not pragmatic
- ▶ **FP is** not suited *to deliver value to the business*

AGENDA

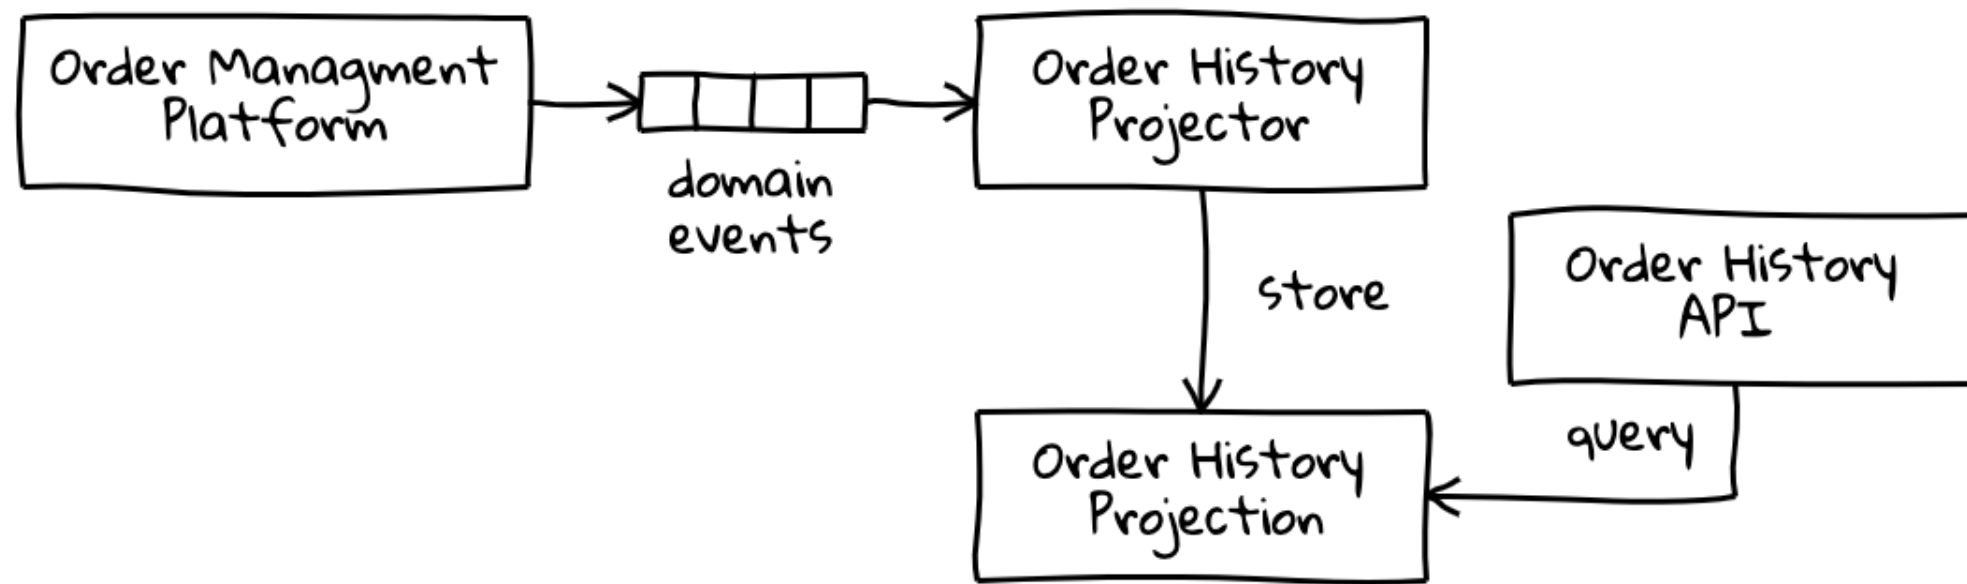
- ▶ A sample architecture
- ▶ Introduce a bunch of building blocks
- ▶ Design architecture components

SAMPLE ARCHITECTURE: *Order History Service*



- ▶ Let's assume we are provided with domain events from an Order Management Platform (e.g. OrderCreated, OrderShipped, etc..), via a RabbitMQ broker
 - ▶ We need to build an Order History Service

ORDER HISTORY SERVICE: *components*



- ▶ a component which projects a (read) model, in a MongoDB collection
- ▶ so that an HTTP service can query the collection returning orders

DISCLAIMER

Our focus here is *NOT* on the System Architecture

We'll just put our attention on implementing an architecture component
(the projector) adopting Functional Programming principles

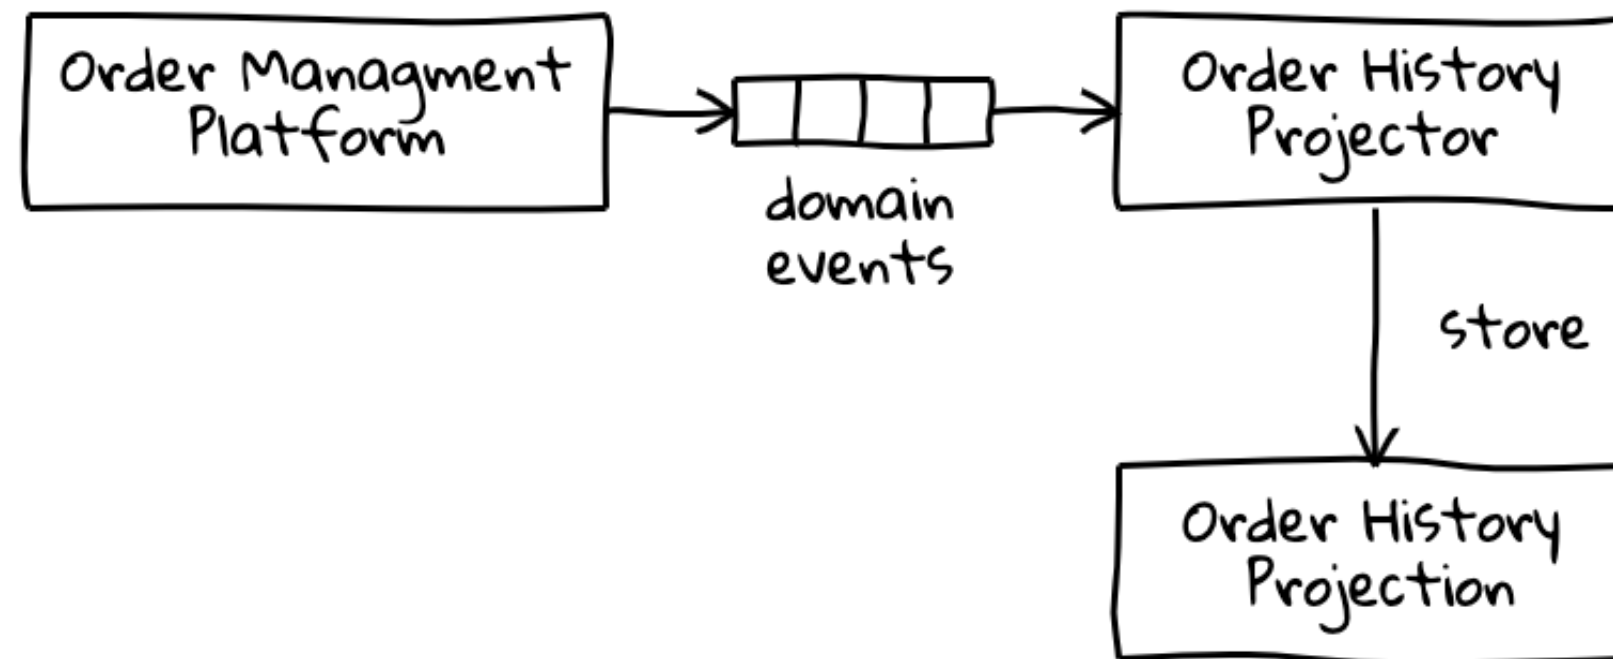
SPOILER

We WON'T be using:

- var
- throw
- methods returning Unit
- poorly typed definitions (Any, Object, etc...)
- low level concurrency mechanisms (Thread, Actor, etc..)

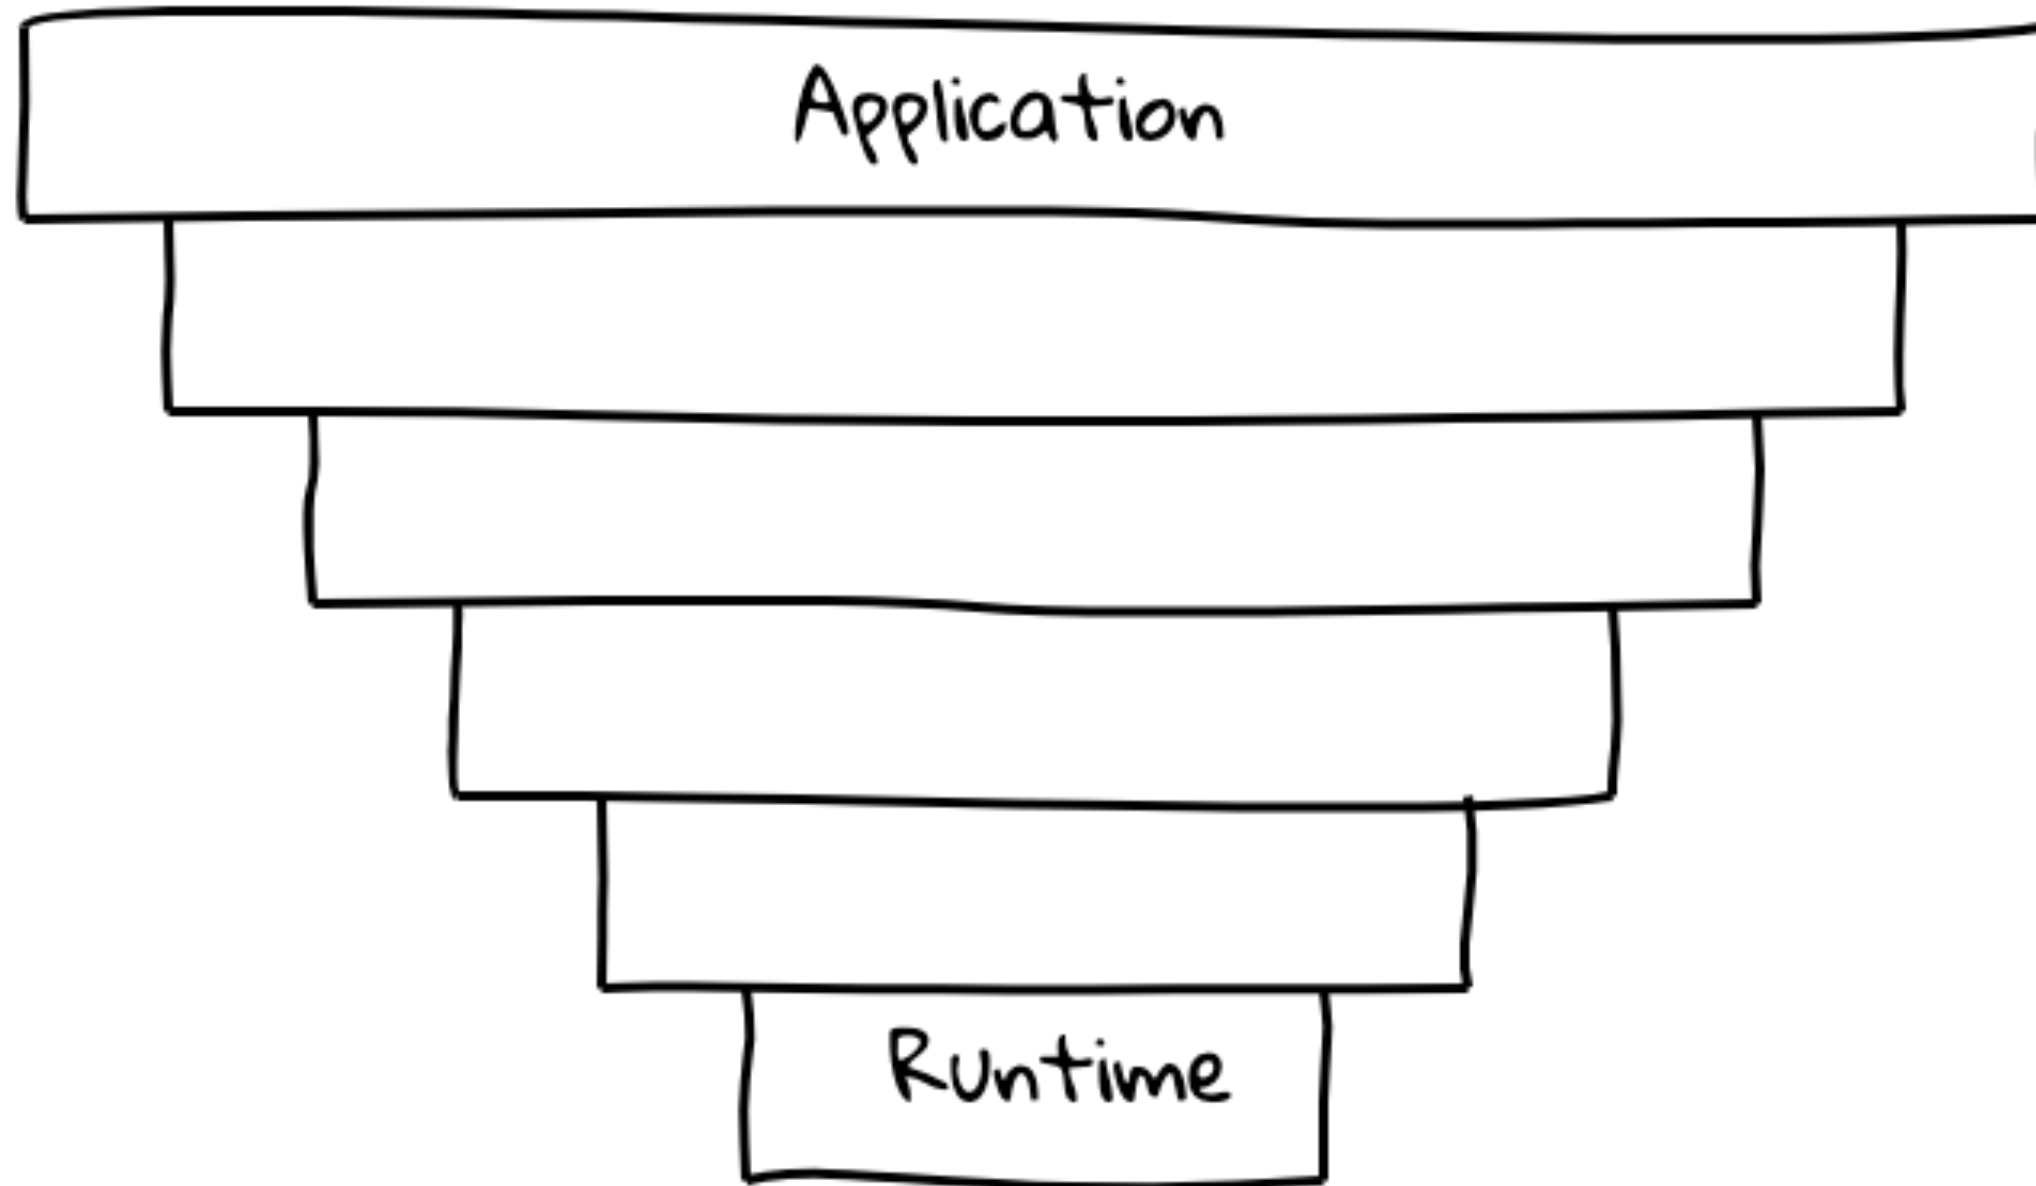
LET'S START

BUILDING A PROJECTOR



- ▶ *Consume* a stream of events from a RabbitMQ queue
- ▶ *Persist* a read model to a MongoDB collection

HOW TO FILL THE ABSTRACTION GAP?



PROJECTOR APPLICATION

1. read a bunch of configs from the env
 2. interact with a RabbitMQ broker
 - 2.1 open a connection
 - 2.2 receive a Stream of events from the given queue
3. interact with a MongoDB cluster
 - 3.1 open a connection
 - 3.2 store the model to the given collection

1. READ A BUNCH OF CONFIGS FROM THE ENV

```
object Mongo {  
  case class Auth(username: String, password: String)  
  case class Config(auth: Auth, addresses: List[String], /*...*/)  
  
  // reading from env variables  
  lazy val config: Config = {  
    val user      = System.getenv("MONGO_USERNAME")  
    val password  = System.getenv("MONGO_PASSWORD")  
    //...reading other env vars ... //  
    Config(Auth(user, password), endpoints, port, db, collection)  
  }  
}
```

**HOW TO TURN AN
API TO BE
FUNCTIONAL™?**

HOW TO TURN AN API TO BE FUNCTIONAL™?

In most cases:

wrap the impure types/operations,
only *expose* a safer version of its operations

CAN FP HELP US

WITH *I/O*

OPERATIONS?

CAN FP HELP US

WITH *side-*

effects?

INTRODUCING IO

A DATA TYPE FOR *encoding effects* AS PURE VALUES

INTRODUCING IO

- ▶ enable capturing and controlling actions - a.k.a effects - that your program wishes to perform within a resource-safe, typed context with seamless support for concurrency and coordination
 - ▶ these effects may be asynchronous (callback-driven) or synchronous (directly returning values); they may return within microseconds or run infinitely.

INTRODUCING IO

A value of type `IO[A]` is a computation that, when evaluated, can perform *effects* before either

- yielding exactly one **result**: a value of type `A`
- raising a **failure** (`Throwable`)

IO VALUES

- ▶ are pure and immutable
- ▶ represents just a description of a side effectful computation
 - ▶ are not evaluated (suspended) until the *end of the world*
 - ▶ respects referential transparency

REFERENTIAL TRANSPARENCY

**An expression may be replaced by its value (or anything having the same value)
without changing the result of the program**

WHY DO WE NEED REFERENTIAL TRANSPARENCY IF IT WORKS ANYWAY?

```
def askInt(): Future[Int] =  
  Future(println("Please, give me a number:"))  
    .flatMap(_ => Future(io.StdIn.readLine().toInt))  
  
def askTwoInt(): Future[(Int, Int)] =  
  for {  
    x <- askInt()  
    y <- askInt()  
  } yield (x , y)  
  
def program(): Future[Unit] =  
  askTwoInt()  
    .flatMap(pair => Future(println(s"Result: ${pair}")))
```

```
> Output:  
> Please, give me a number:  
> 4  
> Please, give me a number:  
> 7  
> Result: (4,7)
```

WHY DO WE NEED REFERENTIAL TRANSPARENCY IF IT WORKS ANYWAY?

```
def askInt(): Future[Int] =  
  Future(println("Please, give me a number:"))  
    .flatMap(_ => Future(io.StdIn.readLine().toInt))  
  
def askTwoInt(): Future[(Int, Int)] =  
  val sameAsk = askInt()  
  for {  
    x <- sameAsk  
    y <- sameAsk  
  } yield (x , y)  
  
def program(): Future[Unit] =  
  askTwoInt()  
    .flatMap(pair => Future(println(s"Result: ${pair}")))
```

```
> Output:  
> Please, give me a number:  
> 4  
> Result: (4,4)
```

**We just wanted to reduce
duplication through an extract var!¹**

¹Example gently stolen from my dear friend <https://github.com/matteobaglini/onion-with-functional-programming>

REFERENTIAL TRANSPARENCY

- ▶ code easier to reason about
 - ▶ code easier to refactor
 - ▶ code easier to compose
- ▶ we're already used to referential transparency since our math lessons!
- ▶ we're already using a lot of data types in a referential transparent manner (e.g. `List`, `Option`, `Try`, `Either`)!

HOW TO TURN AN API TO BE FUNCTIONAL™?

In most cases:

wrap the impure types/operations,
only *expose* a safer version of its operations
which will need to be referential transparent

IO AND COMBINATORS

```
object IO {  
  def delay[A](a: => A): IO[A]  
  def pure[A](a: A): IO[A]  
  def raiseError[A](e: Throwable): IO[A]  
  def sleep(duration: FiniteDuration): IO[Unit]  
  def async[A](k: /* ... */): IO[A]  
  ...  
}
```

```
class IO[A] {  
  def map[B](f: A => B): IO[B]  
  def flatMap[B](f: A => IO[B]): IO[B]  
  def *>[B](fb: IO[B]): IO[B]  
  ...  
}
```

COMPOSING SEQUENTIAL EFFECTS

```
val ioInt: IO[Int] =  
  IO.delay { println("hello") }  
    .map(_ => 1)  
  
val program: IO[Unit] =  
  for {  
    i1 <- ioInt  
    _ <- IO.sleep(i1.second)  
    _ <- IO.raiseError( // not throwing!  
      new RuntimeException("boom!"))  
    i2 <- ioInt //comps is short-circuted  
  } yield ()
```

```
> Output:  
> hello  
> <...1 second...>  
> RuntimeException: boom!
```

**WE ARE
PRACTICAL**

1. READ A BUNCH OF CONFIGS FROM THE ENV, MADE FUNCTIONAL™

```
object Mongo {  
  case class Auth(username: String, password: String)  
  case class Config(auth: Auth, addresses: List[String], /*...*/)  
  
  object Config {  
    // a delayed computation which read from env variables  
    val load: IO[Config] =  
      for {  
        user      <- IO.delay(System.getenv("MONGO_USERNAME"))  
        password <- IO.delay(System.getenv("MONGO_PASSWORD"))  
        //...reading other env vars ... //  
      } yield Config(Auth(user, password), endpoints, port, db, collection)  
  }  
}
```

COMPOSING EFFECTS

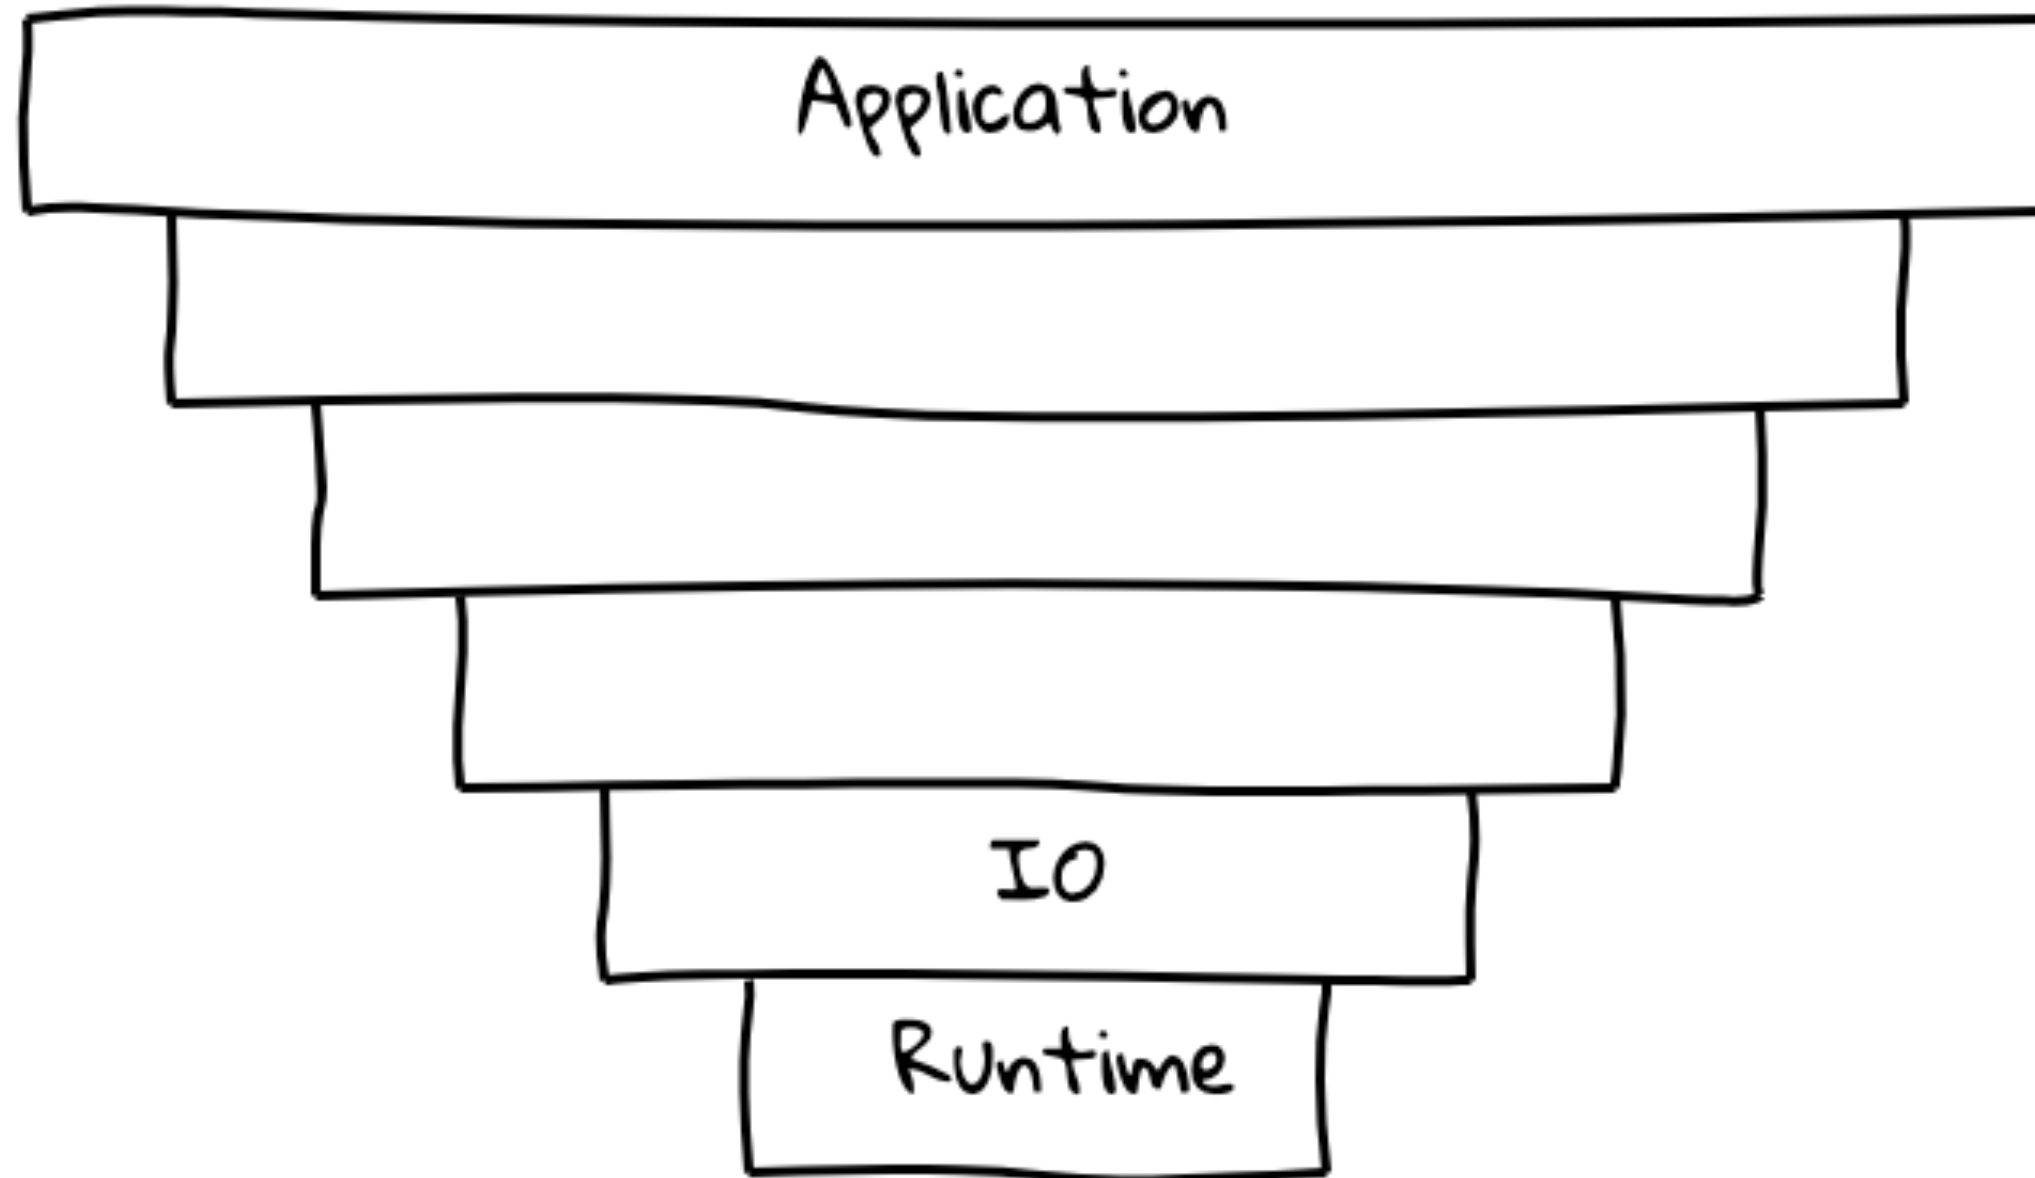
```
val ioOps =  
  for {  
    mongoConfig  <- Mongo.Config.load  
    rabbitConfig <- Rabbit.Config.load  
    // TODO use configs to do something!  
  } yield ()
```

HOW IO VALUES ARE EXECUTED?

If IO values are just a description of **effectful computations** which can be composed and so on...

Who's gonna *run* the suspended computation then?

HOW TO FILL THE ABSTRACTION GAP?



END OF THE WORLD

- ▶ IOApp describes a main which executes an IO
- ▶ as the single entry point to a *pure* program.

```
object OrderHistoryProjectorApp extends IOApp.Simple {  
  override def run: IO[Unit] =  
    for {  
      mongoConfig  <- Mongo.Config.load  
      rabbitConfig <- Rabbit.Config.load  
      // TODO use configs to start the main logic!  
    } yield ()  
}
```

PROJECTOR APPLICATION


1. ~~read a bunch of configs from the env~~
2. interact with a RabbitMQ broker
 - 2.1 open a connection
 - 2.2 receive a Stream of events from the given queue
3. interact with a MongoDB cluster
 - 3.1 open a connection
 - 3.2 store the model to the given collection

2. INTERACT WITH A RABBITMQ BROKER

Using `fs2-rabbit` lib which:

- provides a purely functional api
- let me introduce you a bunch of useful data types

HOW TO FILL THE ABSTRACTION GAP?



2.1. INTERACT WITH A RABBITMQ BROKER

OPEN A CONNECTION

```
val channel: Resource[AMQPChannel] =  
  for {  
    rabbitClient <- RabbitClient.resource(config)  
    channel      <- rabbitClient.createConnectionChannel  
  } yield channel
```

RESOURCE?


INTRODUCING RESOURCE

EFFECTFULLY ALLOCATES AND RELEASES A RESOURCE

EXTREMELY HELPFUL TO WRITE CODE THAT:

- ▶ doesn't leak
- ▶ handles properly terminal signals (e.g. `SIGTERM`) by default (no need to register a shutdown hook)
 - ▶ do the right thing™ by design
- ▶ avoid the need to reboot a container every once in a while :)

HOW TO FILL THE ABSTRACTION GAP?



INTRODUCING RESOURCE

```
object Resource {  
  def make[A](  
    acquire: IO[A])(  
    release: A => IO[Unit]): Resource[A]  
}
```

```
class Resource[A] {  
  def use[B](f: A => IO[B]): IO[B]  
  
  def map[B](f: A => B): Resource[B]  
  def flatMap[B](f: A => Resource[B]): Resource[B]  
  ...  
}
```

NB: not actual code, just a simplification sticking with IO type

MAKING A RESOURCE

```
def mkResource(s: String): Resource[String] = {  
    val acquire =  
        IO.delay(println(s"Acquiring $s")) *> IO.pure(s)  
  
    def release(s: String) =  
        IO.delay(println(s"Releasing $s"))  
  
    Resource.make(acquire)(release)  
}
```

USING A RESOURCE

```
val r: Resource[(String, String)] =  
  for {  
    outer <- mkResource("outer")  
    inner <- mkResource("inner")  
  } yield (outer, inner)  
  
r.use { case (a, b) =>  
  IO.delay(println(s"Using $a and $b"))  
} // IO[Unit]
```

Output:

```
> Acquiring outer  
> Acquiring inner  
> Using outer and inner  
> Releasing inner  
> Releasing outer
```

USING A RESOURCE

```
val sessionPool: Resource[MySessionPool] =  
  for {  
    connection <- openConnection()  
    sessions   <- openSessionPool(connection)  
  } yield sessions  
  
sessionPool.use { sessions =>  
  // use sessions to do whatever things!  
}
```

Output:

```
> Acquiring connection  
> Acquiring sessions  
> Using sessions  
> Releasing sessions  
> Releasing connection
```

GOTCHAS:

- ▶ **Nested resources are released in reverse order of acquisition**
 - ▶ **Easy to lift an `AutoClosable` to `Resource`, via `Resource.fromAutoclosable`**
 - ▶ **Every time you need to use something which implements `AutoClosable`, you should really be using `Resource`!**
- ▶ **You can lift any `IO[A]` into a `Resource[A]` with a no-op release via `Resource.eval`**

WHY NOT SCALA.UTIL.USING?

- ▶ not composable (no `map`, `flatMap`, etc...)
- ▶ no support for properly handling effects

**WE ARE
PRAGMATIC**

2.1. INTERACT WITH A RABBITMQ BROKER

```
object Rabbit {  
  //...  
  def consumerFrom(  
    config: Fs2RabbitConfig,  
    decoder: EnvelopeDecoder[IO, Try[OrderCreatedEvent]]  
  ): Resource[(Acker, Consumer)] =  
    for {  
      rabbitClient <- RabbitClient.resource(config)  
      channel      <- rabbitClient.createConnectionChannel  
      (acker, consumer) <- Resource.eval(  
        rabbitClient.createAckerConsumer(  
          queueName = QueueName("EventsFrom0ms"),  
          channel = channel,  
          decoder = decoder  
        ) // IO[(Acker, Consumer)]  
      )  
    } yield (acker, consumer)  
  
  type Acker      = AckResult => IO[Unit]  
  type Consumer = Stream[AmqpEnvelope[Try[OrderCreatedEvent]]]  
}
```


I HEAR YOU...

```
type Consumer =  
  Stream[AmqpEnvelope[Try[OrderCreatedEvent]]]
```


INTRODUCING STREAM

A SEQUENCE OF EFFECTFUL COMPUTATION

INTRODUCING STREAM

- ▶ *Simplify the way we write concurrent streaming consumers*
- ▶ *Pull-based*, a consumer pulls its values by repeatedly performing pull steps

HOW TO FILL THE ABSTRACTION GAP?



INTRODUCING STREAM

A stream producing output of type `O` and which may evaluate `IO` effects.

```
object Stream {  
  def emit[A](a: A): Stream[A]  
  def emits[A](as: List[A]): Stream[A]  
  def eval[A](f: IO[A]): Stream[A]  
  ...  
}  
  
class Stream[O]{  
  def evalMap[O2](f: O => IO[O2]): Stream[O2]  
  ...  
  def map[O2](f: O => O2): Stream[O2]  
  def flatMap[O2](f: O => Stream[O2]): Stream[O2]  
}
```

NB: not actual code, just a simplification sticking with `IO` type

INTRODUCING STREAM

A sequence of effects...

```
Stream(1,2,3)  
  .repeat  
  .evalMap(i => IO.delay(println(i)))  
  .compile  
  .drain
```

WE DELIVER

CONSUMING A STREAM

```
class OrderHistoryProjector(consumer: Consumer, acker: Acker, logger: Logger) {  
  val project: IO[Unit] =  
    consumer.evalMap { envelope =>  
      envelope.payload match {  
        case Success(event) =>  
          logger.info("Received: " + envelope) *>  
            acker(AckResult.Ack(envelope.deliveryTag))  
        case Failure(e) =>  
          logger.error(e)("Error while decoding") *>  
            acker(AckResult.NAck(envelope.deliveryTag))  
      }  
    }  
  .compile.drain  
}
```


PROJECTOR APPLICATION

1. ~~read a bunch of configs from the env~~
2. ~~interact with a RabbitMQ broker~~
 - 2.1 ~~open a connection~~
 - 2.2 ~~receive a Stream of events from the given queue~~
3. interact with a MongoDB cluster
 - 3.1 open a connection
 - 3.2 store the model to the given collection

3. INTERACT WITH A MONGODB CLUSTER

Using `mongo4cats`, a thin wrapper over the official `mongodb` driver, which exposes purely functional apis

3.1 OPEN A CONNECTION

```
object Mongo {  
  ...  
  def collectionFrom(conf: Config): Resource[MongoCollection[Order]] = {  
    val clientSettings = ??? // conf to mongo-scala-driver settings  
  
    for {  
      client      <- MongoClient.create(settings.build())  
      db          <- Resource.eval(client.getDatabase(conf.databaseName))  
      collection <- Resource.eval(db.getCollectionWithCirceCodecs[Order](conf.collectionName))  
    } yield collection  
  }  
}
```

PROJECTOR APPLICATION

1. ~~read a bunch of configs from the env~~
2. ~~interact with a RabbitMQ broker~~
 - 2.1 ~~open a connection~~
 - 2.2 ~~receive a Stream of events from the given queue~~
3. ~~interact with a MongoDB cluster~~
 - 3.1 ~~open a connection~~
 - 3.2 ~~store the model to the given collection~~

3.2 STORE THE MODEL TO THE GIVEN COLLECTION

```
class EventRepository(collection: MongoClient[Order]) {  
  def store(event: OrderCreatedEvent): IO[Unit] =  
    collection  
      .insertOne(  
        Order(  
          orderNo = OrderNo(event.id),  
          company = Company(event.company),  
          email = Email(event.email),  
          lines = event.lines.map(...)  
        )  
      )  
}
```

3.2 STORE THE MODEL TO THE GIVEN COLLECTION

```
class OrderHistoryProjector(eventRepo: EventRepository,
                           consumer: Consumer,
                           acker: Acker,
                           logger: Logger) {

    val project: IO[Unit] =
        consumer.evalMap { envelope =>
            envelope.payload match {
                case Success(event) =>
                    logger.info("Received: " + envelope) *>
                    eventRepo.store(event) *>
                    acker(AckResult.Ack(envelope.deliveryTag))
                case Failure(e) =>
                    logger.error(e)("Error while decoding") *>
                    acker(AckResult.NAck(envelope.deliveryTag))
            }
        }
        .compile.drain
}
```

PROJECTOR APPLICATION

1. ~~read a bunch of configs from the env~~
2. ~~interact with a RabbitMQ broker~~
 - 2.1 ~~open a connection~~
 - 2.2 ~~receive a Stream of events from the given queue~~
3. ~~interact with a MongoDB cluster~~
 - 3.1 ~~open a connection~~
 - 3.2 ~~store the model to the given collection~~

WIRING

How to achieve separation of concerns and have a good modularity?

WIRING

Constructor Injection!

- ▶ JVM application lifecycle is not so complex
- ▶ *No need for magic*, each dependency can be explicitly injected
- ▶ Acquiring/releasing resources should be handled as an **effect**

INTRODUCING CONSTRUCTOR INJECTION

HOW *not to suffer* WHILE INJECTING DEPENDENCIES

CONSTRUCTOR INJECTION

- ▶ a class with a *private constructor*
- ▶ a companion object with a `fromX/make` method (*smart constructor*)
 1. taking dependencies as input
 2. usually returning `IO/Resource` of the component class

WIRING - CONSTRUCTOR INJECTION

```
class OrderHistoryProjector private (  
  eventRepo: EventRepository,  
  consumer: Consumer,  
  acker: Acker,  
  logger: Logger) {  
  ...  
}  
  
object OrderHistoryProjector {  
  def fromConfigs(mongoConfig: Mongo.Config,  
                 rabbitConfig: Rabbit.Config  
  ): Resource[OrderHistoryProjector] = ...  
}
```

WIRING - CONSTRUCTOR INJECTION

```
object OrderHistoryProjector {  
  def fromConfigs(  
    mongoConfig: Mongo.Config,  
    rabbitConfig: Fs2RabbitConfig  
  ): Resource[OrderHistoryProjector] =  
    for {  
      (ack, consumer) <- Rabbit.consumerFrom(rabbitConfig, eventDecoder)  
      collection      <- Mongo.collectionFrom(mongoConfig)  
      repo            = EventRepository.fromCollection(collection)  
      logger          <- Resource.eval(Slf4jLogger.create)  
    } yield new OrderHistoryProjector(repo, consumer, ack, logger)  
}
```

WIRING - MAIN

```
object OrderHistoryProjectorApp extends IOApp.Simple {  
  
  def run: IO[Unit] =  
    for {  
      mongoConfig  <- Mongo.Config.load  
      rabbitConfig <- Rabbit.Config.load  
  
      _ <- OrderHistoryProjector  
        .fromConfigs(mongoConfig, rabbitConfig) // acquire the needed resources  
        .use(_.project) // start to process the stream of events  
  
    } yield ()  
}
```

ALL DONE!



THERE'S A LOT MORE TO TALK ABOUT

- ▶ How to handle concurrency, execution contexts, blocking ops?
 - ▶ How to track and handle errors?
 - ▶ Do we really need advanced techniques?

CONCLUSIONS

- ▶ a production-ready component in under 300 LOC
- ▶ **only 3 main datatypes:** `IO`, `Resource`, `Stream`
 - ▶ **no variables, no mutable state**
 - ▶ not even the `M` word!
- ▶ *I could have written almost the same code in Kotlin, Swift or.. Haskell!*

REFERENCES

<https://github.com/AL333Z/fp-in-industry>

<https://typelevel.org/cats-effect/>

<https://fs2.io/>

<https://fs2-rabbit.profunktor.dev/>

THANKS

I'VE BEEN LYING TO YOU

STREAM, RESOURCE, RABBITCLIENT AND MONGOCOLLECTION ARE POLYMORPHIC IN THE EFFECT TYPE!

In all the slides I always omitted the additional effect type parameter!

- ▶ `Resource[F, A]`
- ▶ `Stream[F, A]`
- ▶ `RabbitClient[F]`
- ▶ `MongoCollection[F]`

POLYMORPHISM IS GREAT, BUT COMES AT A (LEARNING) COST!